



INTER IIT TECH MEET 14.0

Client Centric WiFi RRM

INTER IIT TECH MEET 14.0

Detailed Design Document - Causal Inference Engine

Team 33

ARISTA

Contents

1 System Overview	3
1.1 Key Objectives	3
1.2 Architecture Overview	3
2 Data Structures	4
2.1 Configuration Structure	4
2.2 Result Structure	4
2.2.1 Primary Outputs	4
2.2.2 Validation Metrics	4
2.2.3 Method-Specific Estimates	5
2.2.4 Data Quality Metrics	5
2.2.5 Model Information	5
3 Input Data Specification	6
3.1 Action File Format	6
3.2 Graph Snapshot Files	6
3.2.1 Directory Structure	6
3.2.2 Node Features Schema	6
3.2.3 Edge Features Schema	7
4 Problem Formulation	8
4.1 Causal Estimand	8
4.2 Identification Assumptions	8
5 QoE Computation Algorithm	9
5.1 QoE Formula	9
5.2 Retry Rate Estimation	9
5.3 Latency Estimation	9
5.4 Algorithm: Compute QoE per AP	10
6 Method 1: Propensity Score Matching (PSM)	11
6.1 Theory	11
6.2 Algorithm: Propensity Score Matching	11
6.3 PSM Process Flow	12
6.4 Implementation Details	12
7 Method 2: Double/Debiased Machine Learning (DML)	13
7.1 Theory	13
7.2 Cross-Fitting to Avoid Overfitting	13
7.3 Algorithm: Double Machine Learning	14
7.4 Implementation Details	14
8 Aggregation and Statistical Inference	15
8.1 Meta-Estimator	15
8.2 Confidence Interval Construction	15
8.3 Hypothesis Testing	15
9 Covariate Balance Assessment	16
9.1 Standardized Mean Difference (SMD)	16
9.2 Overall Balance Score	16
10 Reward Adjustment for RL Integration	17
10.1 Causal Reward Function	17
10.2 RL-Compatible Reward Computation	17
10.3 Rationale for Reward Shaping	17
11 Model Persistence and Progressive Learning	18
11.1 Strategy	18
11.2 Cached Artifacts	18
11.3 Progressive Learning Workflow	18

12 Integration with RL System	19
12.1 Usage Example	19
12.2 Batch Integration	19
13 Computational Complexity	20
14 Summary and Conclusion	21
14.1 Key Achievements	21
14.2 System Capabilities	21
14.3 Integration Benefits	21
14.4 Future Enhancements	21

1 System Overview

The Causal Reward Engine is designed to validate Radio Resource Management (RRM) actions by computing causally-validated rewards for Reinforcement Learning (RL) agents. Unlike traditional correlation-based metrics, this system employs rigorous causal inference methods to isolate the true effect of RRM interventions on network Quality of Experience (QoE).

1.1 Key Objectives

- Replace **GNN predictions** with real causal rewards during deployment
- Validate **counterfactual reasoning** by comparing treatment vs control groups
- Handle **small sample sizes** through model persistence and progressive learning
- Maintain **RL compatibility** by matching existing reward structures
- Provide **statistical guarantees** through significance testing and confidence intervals

1.2 Architecture Overview

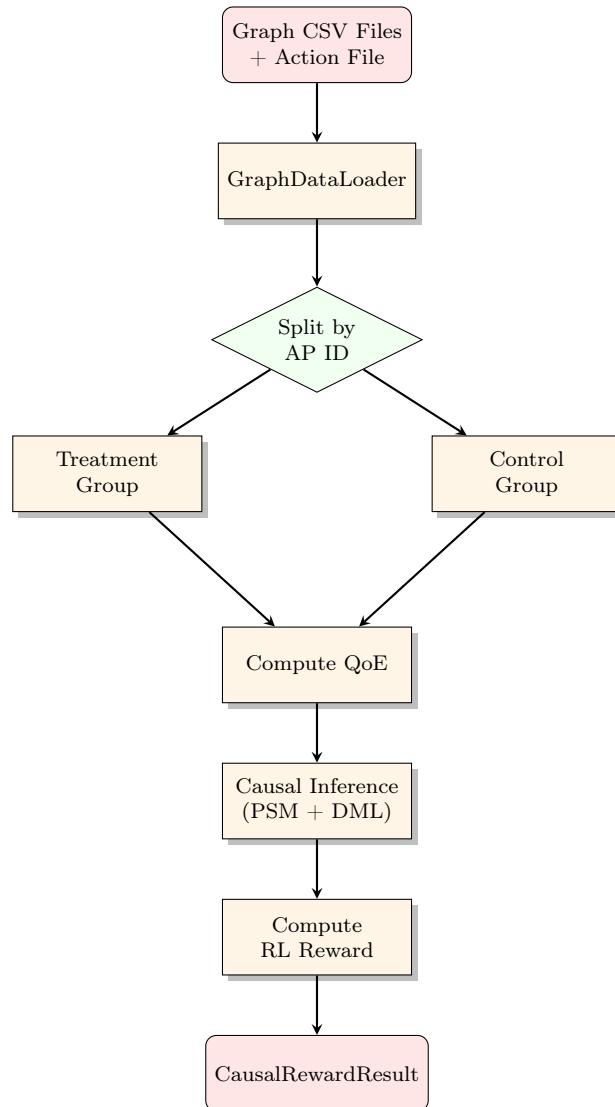


Figure 1: Causal Reward Engine Architecture

2 Data Structures

2.1 Configuration Structure

The **CausalConfig** class encapsulates all configuration parameters for the causal inference engine.

Table 1: CausalConfig Field Specifications

Field Name	Type	Default	Description
graph_data_dir	str	-	Directory containing graph CSV files
action_file_path	str	-	Path to action file (apply/skip labels)
min_graphs_required	int	10	Minimum graphs for valid analysis
max_graphs_to_load	int	20	Maximum graphs to process
model_cache_dir	str	-	Directory for model persistence
enable_model_persistence	bool	True	Enable progressive learning
min_samples_for_training	int	30	Threshold for model training
psm_caliper	float	0.15	Maximum propensity score distance
psm_k_neighbors	int	3	Number of nearest neighbors for matching
significance_level	float	0.05	Statistical significance threshold
min_sample_size_per_group	int	5	Minimum samples per group
w_penalty	float	3.0	Weight for QoE decrease (RL agent)
w_uplift	float	1.0	Weight for QoE increase (RL agent)
w_state	float	0.5	State bias weight (RL agent)
reward_clip_min	float	-10.0	Minimum reward value
reward_clip_max	float	10.0	Maximum reward value
retry_penalty_weight	float	10.0	QoE penalty for retries
latency_penalty_weight	float	0.1	QoE penalty for latency

2.2 Result Structure

The **CausalRewardResult** class contains the complete output of causal reward computation, organized into four categories.

2.2.1 Primary Outputs

Table 2: CausalRewardResult - Primary Outputs

Field Name	Type	Description
causal_reward	float	Final reward for RL agent (clipped to $[-10, 10]$)
causal_effect	float	Estimated treatment effect on QoE

2.2.2 Validation Metrics

Table 3: CausalRewardResult - Validation Metrics

Field Name	Type	Description
is_significant	bool	Statistical significance ($p < 0.05$)
p_value	float	Two-tailed p-value from t-test
confidence_interval	(float, float)	95% CI for effect

2.2.3 Method-Specific Estimates

Table 4: CausalRewardResult - Method-Specific Estimates

Field Name	Type	Description
psm_effect	Optional[float]	Effect from PSM method
dml_effect	Optional[float]	Effect from DML method

2.2.4 Data Quality Metrics

Table 5: CausalRewardResult - Data Quality Metrics

Field Name	Type	Description
n_treatment	int	Number of treatment units
n_control	int	Number of control units
balance_score	float	Covariate balance quality $\in [0, 1]$
qoe_treatment_mean	float	Average QoE in treatment group
qoe_control_mean	float	Average QoE in control group
observed_delta	float	Raw difference (before causal adjustment)

2.2.5 Model Information

Table 6: CausalRewardResult - Model Information

Field Name	Type	Description
used_cached_models	bool	Whether cached models were used
model_sample_count	int	Number of samples in cached models

3 Input Data Specification

3.1 Action File Format

The action file determines treatment/control group assignment.

Table 7: Action File Schema (`approved_changes.csv`)

Column	Type	Description
<code>ap_id</code>	int	Unique Access Point identifier
<code>channel_width</code>	int	Channel width in MHz (e.g., 20, 40, 80)
<code>tx_power</code>	float	Transmit power in dBm
<code>obss_pd</code>	float	OBSS-PD threshold in dBm
<code>action</code>	str	'apply' (treatment) or 'skip' (control)

Example:

Listing 1: Sample Action File

```

1 ap_id,channel_width,tx_power,obss_pd,action
2 4,80,3,-52,apply
3 6,20,18,-52,apply
4 2,80,3,-52,apply
5 5,20,23,-82,skip
6 1,20,23,-82,skip

```

3.2 Graph Snapshot Files

Graph snapshots capture network state at multiple timestamps.

3.2.1 Directory Structure

Listing 2: Expected Directory Layout

```

1 causal_data/graphs/
2   interference_graph_nodes/
3     interference_graph_nodes_t60.0s.csv
4     interference_graph_nodes_t150.0s.csv
5   ...
6   interference_graph_edges/
7     interference_graph_edges_t60.0s.csv
8     interference_graph_edges_t150.0s.csv
9   ...

```

3.2.2 Node Features Schema

Table 8: Node CSV Schema (AP-level features)

Column	Type	Description
<code>ap_id</code>	int	Access Point identifier
<code>channel</code>	int	Operating channel number
<code>power_dbm</code>	float	Transmit power in dBm
<code>width_mhz</code>	int	Channel width in MHz
<code>obss_pd_dbm</code>	float	OBSS-PD threshold
<code>client_count</code>	int	Number of associated clients
<code>airtime_util</code>	float	Airtime utilization $\in [0, 1]$
<code>tx_bytes</code>	int	Transmitted bytes
<code>rx_bytes</code>	int	Received bytes
<code>avg_throughput_mbps</code>	float	Average throughput in Mbps
<code>avg_client_rssi</code>	float	Average client RSSI in dBm
<code>edge_client_count</code>	int	Number of edge clients (low RSSI)
<code>client_rssi_p90</code>	float	90th percentile client RSSI
<code>client_rssi_p10</code>	float	10th percentile client RSSI

3.2.3 Edge Features Schema

Table 9: Edge CSV Schema (Inter-AP relationships)

Column	Type	Description
source_ap_id	int	Source AP identifier
dest_ap_id	int	Destination AP identifier
coupling_dbm	float	Signal coupling strength
channel_overlap	float	Channel overlap fraction $\in [0, 1]$
interference_power_dbm	float	Interference power in dBm
is_cochannel	bool	True if same channel
channel_separation	int	Channel number difference
roaming_client_count	int	Clients roaming between APs
spectral_overlap	float	Spectral overlap metric
interference_power_linear	float	Linear interference power
hidden_node_proxy	float	Hidden node indicator
power_delta_dbm	float	Power difference between APs

4 Problem Formulation

To optimize WiFi Radio Resource Management (RRM), it is insufficient to simply observe correlations between actions (e.g., changing channel width) and network metrics (e.g., throughput). Observational data is often confounded by environmental factors such as user density or interference.

This module formally defines the causal inference framework used to isolate the true effect of RRM interventions. Our goal is to estimate the counterfactual: *what would have happened to the network performance had the RRM action not been taken?*

4.1 Causal Estimand

We define the fundamental variables required to estimate the causal effect. In our context, a "unit" i represents a specific Access Point (AP) or a client session.

Notation:

- Y_{1i} : Potential outcome if unit i receives treatment (e.g., new power level)
- Y_{0i} : Potential outcome if unit i receives control (e.g., existing power level)
- $T_i \in \{0, 1\}$: Treatment indicator
- X_i : Vector of observed covariates (e.g., RSSI, noise floor, client count)
- $Y_i = T_i Y_{1i} + (1 - T_i) Y_{0i}$: Observed outcome

Our primary objective is to estimate the effect on those units that actually received the intervention:

$$\tau_{ATT} = E[Y_{1i} - Y_{0i} \mid T_i = 1] \quad (1)$$

4.2 Identification Assumptions

Causal inference from observational data requires strict assumptions to ensure that the estimated effect is not driven by selection bias. We invoke the following three standard assumptions:

1. Unconfoundedness/Ignorability:

$$(Y_{1i}, Y_{0i}) \perp\!\!\!\perp T_i \mid X_i$$

Treatment assignment is independent of potential outcomes conditional on covariates. This implies we have measured all relevant confounders affecting both RRM decisions and network performance.

2. Common Support/Overlap:

$$0 < P(T_i = 1 \mid X_i) < 1 \quad \forall X_i$$

Both treatment and control units exist at all covariate values, ensuring valid comparisons can be made across the state space.

3. SUTVA (Stable Unit Treatment Value Assumption):

- No interference between units (an action on AP A does not implicitly affect AP B)
- Single version of treatment

5 QoE Computation Algorithm

Before performing causal inference, we must compute a Quality of Experience (QoE) metric for each AP. This metric combines throughput with penalties for retries and latency, matching the RL agent's reward structure.

5.1 QoE Formula

$$\text{QoE}_i = \log(T_i + 1) - w_r \cdot R_i - w_\ell \cdot L_i \quad (2)$$

where:

- T_i : Average throughput (Mbps) for AP i
- R_i : Retry rate estimate for AP i
- L_i : Latency estimate (ms) for AP i
- $w_r = 10.0$: Retry penalty weight
- $w_\ell = 0.1$: Latency penalty weight

5.2 Retry Rate Estimation

Since graph CSV files do not directly contain retry rate, we estimate it from interference and airtime utilization:

$$R_i = 0.02 + 0.08 \cdot \frac{I_i + 90}{30} \cdot A_i \quad (3)$$

where:

- I_i : Interference power (dBm), normalized from $[-90, -60]$ range
- A_i : Airtime utilization $\in [0, 1]$

Rationale: High interference combined with high airtime leads to more packet collisions and retransmissions.

5.3 Latency Estimation

$$L_i = 10 + 20 \cdot A_i + 0.5 \cdot C_i \quad (4)$$

where:

- A_i : Airtime utilization $\in [0, 1]$
- C_i : Client count

Rationale: High airtime and many clients increase queueing delays and contention.

5.4 Algorithm: Compute QoE per AP

Algorithm 1 Compute QoE for All APs

Input: DataFrame \mathcal{D} with columns: `ap_id`, `avg_throughput_mbps`, `client_count`, `airtime_util`, `interference_power_dbm`

Output: DataFrame \mathcal{D} with additional column: `qoe_metric`

```

1:  $w_r \leftarrow 10.0$                                 ▷ Retry penalty weight
2:  $w_\ell \leftarrow 0.1$                                 ▷ Latency penalty weight
3: for each row  $i$  in  $\mathcal{D}$  do
4:    $T_i \leftarrow \mathcal{D}[i].\text{avg\_throughput\_mbps}$ 
5:    $A_i \leftarrow \mathcal{D}[i].\text{airtime\_util}$ 
6:    $C_i \leftarrow \mathcal{D}[i].\text{client\_count}$ 
7:    $I_i \leftarrow \mathcal{D}[i].\text{interference\_power\_dbm}$           ▷ Fill NaN with  $-80$ 
8:    $I_{\text{norm}} \leftarrow (I_i + 90)/30$                       ▷ Estimate retry rate
9:    $R_i \leftarrow 0.02 + 0.08 \cdot I_{\text{norm}} \cdot A_i$         ▷ Normalize to  $[0, 1]$ 
10:   $L_i \leftarrow 10 + 20 \cdot A_i + 0.5 \cdot C_i$             ▷ Estimate latency
11:   $\text{QoE}_i \leftarrow \log(T_i + 1) - w_r \cdot R_i - w_\ell \cdot L_i$  ▷ Compute QoE
12:   $\mathcal{D}[i].\text{qoe\_metric} \leftarrow \text{QoE}_i$ 
13: end for
14: return  $\mathcal{D}$ 

```

6 Method 1: Propensity Score Matching (PSM)

Propensity Score Matching creates a "quasi-randomized" experiment from observational data. By pairing treated APs with control APs that have similar probabilities of receiving the treatment, we reduce selection bias caused by the inherent policy of the legacy RRM system.

6.1 Theory

The **propensity score** is the probability of treatment given covariates. It serves as a dimensionality reduction technique, collapsing multidimensional covariates X into a single scalar.

$$e(X_i) = P(T_i = 1 \mid X_i) \quad (5)$$

Rosenbaum-Rubin Theorem: If unconfoundedness holds given X , it also holds given $e(X)$:

$$(Y_{1i}, Y_{0i}) \perp T_i \mid e(X_i)$$

6.2 Algorithm: Propensity Score Matching

Algorithm 2 Propensity Score Matching (PSM)

Input: Treatment data \mathcal{D}_1 , Control data \mathcal{D}_0 , Covariates $\mathcal{X} = \{X_1, \dots, X_p\}$

Input: Configuration: K (neighbors), δ (caliper threshold)

Output: ATT estimate $\hat{\tau}_{PSM}$

```

1: Step 1: Prepare combined dataset
2:  $\mathcal{D}_1.\text{treatment} \leftarrow 1$ 
3:  $\mathcal{D}_0.\text{treatment} \leftarrow 0$ 
4:  $\mathcal{D} \leftarrow \text{concatenate}(\mathcal{D}_1, \mathcal{D}_0)$ 
5: Step 2: Extract and scale covariates
6:  $X \leftarrow \mathcal{D}[\mathcal{X}]$                                 ▷ Extract covariate matrix
7:  $X \leftarrow \text{fillna}(X, 0)$                       ▷ Impute missing values
8:  $X_{\text{scaled}} \leftarrow \text{StandardScaler.fit\_transform}(X)$ 
9: Step 3: Estimate propensity scores
10: Train logistic regression:  $\hat{e}(X) = \sigma(\beta_0 + \sum_{j=1}^p \beta_j X_j)$ 
11:  $\hat{e}_i \leftarrow \text{LogisticRegression.predict_proba}(X_{\text{scaled}}, i)[:, 1]$  for all  $i$ 
12:  $\mathcal{D}.\text{propensity\_score} \leftarrow \hat{e}$ 
13: Step 4: Perform nearest neighbor matching
14:  $\mathcal{T} \leftarrow \{i : \mathcal{D}_i.\text{treatment} = 1\}$           ▷ Treatment indices
15:  $\mathcal{C} \leftarrow \{i : \mathcal{D}_i.\text{treatment} = 0\}$           ▷ Control indices
16: Initialize  $K$ -NN model on control propensity scores
17:  $\text{effects} \leftarrow []$                                 ▷ Store matched pair effects
18: for each  $i \in \mathcal{T}$  do
19:   Find  $K$  nearest neighbors  $\mathcal{N}_i \subset \mathcal{C}$  where:
20:      $|\hat{e}_i - \hat{e}_j| \leq \delta$  for all  $j \in \mathcal{N}_i$            ▷ Caliper constraint
21:     if  $\mathcal{N}_i \neq \emptyset$  then
22:        $Y_i^{\text{treat}} \leftarrow \mathcal{D}_i.\text{qoe\_metric}$ 
23:        $Y_i^{\text{control}} \leftarrow \frac{1}{|\mathcal{N}_i|} \sum_{j \in \mathcal{N}_i} \mathcal{D}_j.\text{qoe\_metric}$ 
24:        $\text{effects.append}(Y_i^{\text{treat}} - Y_i^{\text{control}})$ 
25:     end if
26:   end for
27: Step 5: Compute ATT
28:  $\hat{\tau}_{PSM} \leftarrow \frac{1}{|\text{effects}|} \sum_{\text{effect} \in \text{effects}} \text{effect}$ 
29: return  $\hat{\tau}_{PSM}$ 

```

6.3 PSM Process Flow

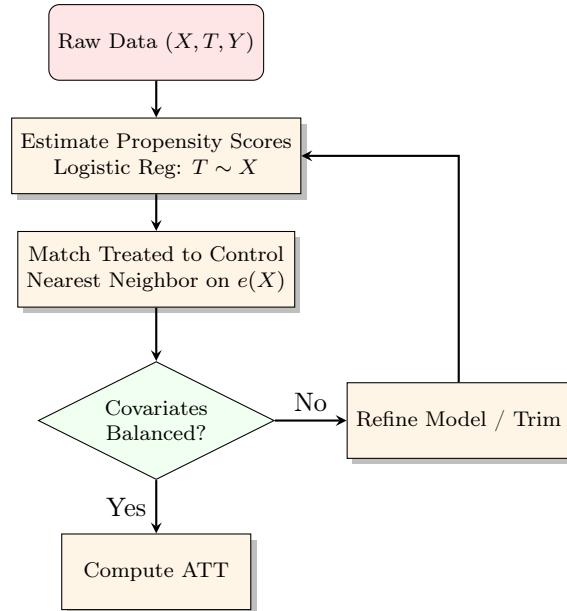


Figure 2: Propensity Score Matching Workflow

6.4 Implementation Details

To ensure robustness in the WiFi environment where noise is high:

- **Model:** Logistic regression with L2 regularization (max iterations = 1000) to handle collinear network metrics
- **Matching:** K-nearest neighbors ($K = 3$ default) to reduce variance
- **Caliper:** $\delta = 0.15$ standard deviations of propensity score
- **Variance estimation:** Bootstrap with replacement ($B = 100$) to generate confidence intervals

7 Method 2: Double/Debiased Machine Learning (DML)

Modern WiFi networks generate high-dimensional telemetry data. Traditional linear models may fail to capture complex non-linear relationships, while standard Machine Learning models introduce regularization bias. Double Machine Learning (DML) solves this by using ML to model nuisance parameters while preserving valid statistical inference for the treatment effect.

7.1 Theory

DML uses two ML models: one to predict the outcome from covariates, and another to predict the treatment from covariates. It then correlates the residuals of these predictions.

$$\psi(W; \tau, \eta) = (Y - \ell(X)) - \tau(T - m(X)) \quad (6)$$

where $m(X) = E[T | X]$ (propensity score) and $\ell(X) = E[Y | X]$ (outcome regression).

The key insight is that this score is **orthogonal** to nuisance parameters $\eta = (m, \ell)$, meaning small errors in estimating m and ℓ do not bias the treatment effect τ .

7.2 Cross-Fitting to Avoid Overfitting

To prevent overfitting bias from learning nuisance functions and treatment effect on same data, DML employs **K-fold cross-fitting**.

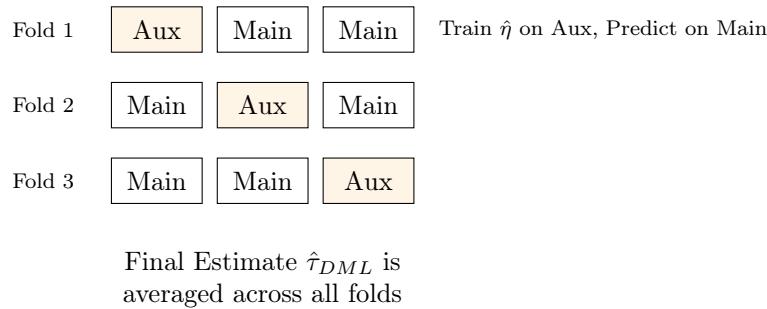


Figure 3: K-Fold Cross-Fitting Concept (K=3)

7.3 Algorithm: Double Machine Learning

Algorithm 3 Double/Debiased Machine Learning (DML)

Input: Treatment data \mathcal{D}_1 , Control data \mathcal{D}_0 , Covariates \mathcal{X}

Input: Configuration: K (number of folds)

Output: ATT estimate $\hat{\tau}_{DML}$

```

1: Step 1: Prepare combined dataset
2:  $\mathcal{D}_1.\text{treatment} \leftarrow 1$ 
3:  $\mathcal{D}_0.\text{treatment} \leftarrow 0$ 
4:  $\mathcal{D} \leftarrow \text{concatenate}(\mathcal{D}_1, \mathcal{D}_0)$ 
5:  $\mathcal{D} \leftarrow \text{shuffle}(\mathcal{D})$                                 ▷ Random shuffle for cross-validation
6: Step 2: Extract features
7:  $X \leftarrow \mathcal{D}[\mathcal{X}].\text{fillna}(0).\text{values}$ 
8:  $T \leftarrow \mathcal{D}.\text{treatment}.\text{values}$ 
9:  $Y \leftarrow \mathcal{D}.\text{qoe\_metric}.\text{values}$ 
10:  $n \leftarrow \text{len}(\mathcal{D})$ 
11: Step 3: K-Fold Cross-Fitting
12:  $\theta_{\text{estimates}} \leftarrow []$                                 ▷ Store treatment effects from each fold
13: for  $k = 1$  to  $K$  do
14:   end if                                                 ▷ Split data into auxiliary and main samples
15:    $\mathcal{I}_{\text{aux}} \leftarrow \{\text{fold } k \text{ indices}\}$ 
16:    $\mathcal{I}_{\text{main}} \leftarrow \{1, \dots, n\} \setminus \mathcal{I}_{\text{aux}}$ 
17:    $X_{\text{aux}}, T_{\text{aux}}, Y_{\text{aux}} \leftarrow X[\mathcal{I}_{\text{aux}}], T[\mathcal{I}_{\text{aux}}], Y[\mathcal{I}_{\text{aux}}]$ 
18:    $X_{\text{main}}, T_{\text{main}}, Y_{\text{main}} \leftarrow X[\mathcal{I}_{\text{main}}], T[\mathcal{I}_{\text{main}}], Y[\mathcal{I}_{\text{main}}]$ 
19:   if  $|\text{unique}(T_{\text{aux}})| < 2$  then
20:     continue                                              ▷ Skip if insufficient treatment variation
21:   end if
22:   Train:  $\hat{m}_k(X) = P(T = 1 | X)$  using LogisticRegression on  $(X_{\text{aux}}, T_{\text{aux}})$            ▷ Fit propensity model on auxiliary sample
23:    $\hat{m}_{\text{main}} \leftarrow \hat{m}_k.\text{predict\_proba}(X_{\text{main}})[:, 1]$ 
24:   Train:  $\hat{\ell}_k(X) = E[Y | X]$  using GradientBoostingRegressor on  $(X_{\text{aux}}, Y_{\text{aux}})$            ▷ Fit outcome model on auxiliary sample
25:    $\hat{\ell}_{\text{main}} \leftarrow \hat{\ell}_k.\text{predict}(X_{\text{main}})$ 
26:   Compute orthogonal residuals on main sample                                         ▷ Compute orthogonal residuals on main sample
27:    $V \leftarrow T_{\text{main}} - \hat{m}_{\text{main}}$                                               ▷ Treatment residuals
28:    $W \leftarrow Y_{\text{main}} - \hat{\ell}_{\text{main}}$                                               ▷ Outcome residuals
29:   Compute fold-specific treatment effect                                         ▷ Compute fold-specific treatment effect
30:   denominator  $\leftarrow \sum_i V_i \cdot T_{\text{main}, i}$ 
31:   if  $|\text{denominator}| > 10^{-6}$  then
32:      $\hat{\theta}_k \leftarrow \frac{\sum_i V_i \cdot W_i}{\text{denominator}}$ 
33:     theta_estimates.append( $\hat{\theta}_k$ )
34:   end if
35:   end for
36: Step 4: Aggregate across folds
37: if theta_estimates  $= \emptyset$  then
38:   return  $\bar{Y}_1 - \bar{Y}_0$                                               ▷ Fallback to simple difference
39: else
40:    $\hat{\tau}_{DML} \leftarrow \frac{1}{K} \sum_{k=1}^K \hat{\theta}_k$ 
41:   return  $\hat{\tau}_{DML}$ 
42: end if

```

7.4 Implementation Details

- **Propensity Model:** Logistic Regression (max_iter=500, L2 regularization)
- **Outcome Model:** Gradient Boosting Regressor (n_estimators=50, learning_rate=0.1)
- **Folds:** $K = 3$ for computational efficiency with small samples
- **Robustness:** Skips folds with insufficient treatment variation or numerical instability

8 Aggregation and Statistical Inference

Each method (PSM, DML) has different strengths and biases. To produce a robust causal estimate, we employ a meta-aggregation strategy that combines estimates while quantifying uncertainty.

8.1 Meta-Estimator

$$\hat{\tau}_{\text{causal}} = \text{median}\{\hat{\tau}_{PSM}, \hat{\tau}_{DML}\} \quad (7)$$

Rationale: The median is robust to outliers from model misspecification, whereas the mean would be sensitive to a single failed method.

8.2 Confidence Interval Construction

We derive uncertainty from cross-method variance:

Standard Error:

$$\text{SE}(\hat{\tau}) = \frac{1}{\sqrt{M}} \sqrt{\frac{1}{M-1} \sum_{m=1}^M (\hat{\tau}_m - \hat{\tau})^2} \quad (8)$$

where M is the number of successful methods.

95% Confidence Interval:

$$\text{CI}_{0.95} = [\hat{\tau} - 1.96 \times \text{SE}(\hat{\tau}), \hat{\tau} + 1.96 \times \text{SE}(\hat{\tau})] \quad (9)$$

8.3 Hypothesis Testing

Null Hypothesis: $H_0 : \tau = 0$ (no treatment effect)

Test Statistic:

$$t = \frac{\hat{\tau}}{\text{SE}(\hat{\tau}) + \epsilon} \quad (10)$$

where $\epsilon = 10^{-10}$ prevents division by zero.

P-value:

$$p = 2 \times (1 - \Phi(|t|)) \quad (11)$$

where Φ is the Student's t-distribution CDF with $df = M - 1$.

Decision Rule:

- If $p < \alpha$ (typically $\alpha = 0.05$): Reject H_0 , effect is **significant**
- If $p \geq \alpha$: Fail to reject H_0 , effect is **not significant**

9 Covariate Balance Assessment

Before accepting a causal estimate, we validate the quality of matching/weighting by checking covariate balance between treatment and control groups.

9.1 Standardized Mean Difference (SMD)

For each covariate X_j :

$$\boxed{\text{SMD}_j = \frac{\bar{X}_{j,1} - \bar{X}_{j,0}}{\sqrt{(s_{j,1}^2 + s_{j,0}^2)/2}}} \quad (12)$$

where:

- $\bar{X}_{j,1}$: Mean of covariate j in treatment group
- $\bar{X}_{j,0}$: Mean of covariate j in control group
- $s_{j,1}^2, s_{j,0}^2$: Variances in treatment and control groups

Balance Criterion: $|\text{SMD}_j| < 0.1$ indicates excellent balance.

9.2 Overall Balance Score

We compute a global balance score $\beta \in [0, 1]$:

$$\boxed{\beta = \max \left(0, 1 - \frac{\overline{\text{SMD}}}{0.5} \right)} \quad (13)$$

where $\overline{\text{SMD}} = \frac{1}{p} \sum_{j=1}^p |\text{SMD}_j|$ is the average absolute SMD across all covariates.

Interpretation:

- $\beta = 1.0$: Perfect balance ($\overline{\text{SMD}} = 0$)
- $\beta \geq 0.8$: Good balance
- $\beta < 0.5$: Poor balance, estimates may be unreliable

10 Reward Adjustment for RL Integration

This section connects the Causal Inference module to the Reinforcement Learning agent. Standard RL rewards can be noisy; by filtering through causal validation, we ensure the agent learns only from statistically significant improvements.

10.1 Causal Reward Function

The final reward incorporates both the causal effect and its statistical significance:

$$R_{\text{causal}} = \begin{cases} R_{\text{RL}}(\hat{\tau}) \times \beta & \text{if } p < \alpha \\ R_{\text{RL}}(\hat{\tau}) \times 0.1 & \text{if } p \geq \alpha \end{cases} \quad (14)$$

where $R_{\text{RL}}(\cdot)$ is the RL-compatible reward function (detailed below), and β is the balance score.

10.2 RL-Compatible Reward Computation

Algorithm 4 Compute RL-Compatible Reward

Input: Causal effect $\hat{\tau}$, Treatment data \mathcal{D}_1 , Is significant: sig

Input: Weights: $w_{\text{penalty}} = 3.0$, $w_{\text{uplift}} = 1.0$, $w_{\text{state}} = 0.5$

Output: Causal reward $R_{\text{causal}} \in [-10, 10]$

```

1: Step 1: Apply significance discount
2:  $\Delta_{\text{QoE}} \leftarrow \hat{\tau}$ 
3: if sig = False then
4:    $\Delta_{\text{QoE}} \leftarrow \Delta_{\text{QoE}} \times 0.1$                                 ▷ 90% discount for non-significant
5: end if
6: Step 2: Apply asymmetric weights
7: if  $\Delta_{\text{QoE}} < 0$  then
8:    $R_{\Delta} \leftarrow w_{\text{penalty}} \times \Delta_{\text{QoE}}$                             ▷ Penalize degradation 3x
9: else
10:   $R_{\Delta} \leftarrow w_{\text{uplift}} \times \Delta_{\text{QoE}}$                                 ▷ Reward improvement 1x
11: end if
12: Step 3: Add state bias
13:  $\text{QoE}_{\text{next}} \leftarrow \text{mean}(\mathcal{D}_1.\text{qoe\_metric})$ 
14:  $R \leftarrow R_{\Delta} + w_{\text{state}} \times \text{QoE}_{\text{next}}$ 
15: Step 4: Clip to bounds
16:  $R_{\text{causal}} \leftarrow \text{clip}(R, -10, 10)$ 
17: return  $R_{\text{causal}}$ 

```

10.3 Rationale for Reward Shaping

- **Significant effect + Good balance:** Full reward scaled by balance score β . This encourages actions where causal link is clear
- **Non-significant effect:** 90% discount prevents "hallucinated" learning from noise
- **Asymmetric penalties:** Penalize QoE degradation 3x more than rewarding improvement (conservative RL policy)
- **State bias:** Encourages maintaining high absolute QoE, not just relative improvements

11 Model Persistence and Progressive Learning

Real-world WiFi deployments may have limited initial data. To handle this, the engine implements **progressive learning** through model caching.

11.1 Strategy

1. **Initial Phase** ($n < 30$ samples): Train models from scratch, cache for future use
2. **Growth Phase** ($30 \leq n < 100$ samples):
 - If new batch has ≥ 30 samples: Retrain and update cache
 - If new batch has < 30 samples: Use cached models (avoid overfitting)
3. **Mature Phase** ($n \geq 100$ samples): Always retrain on latest data

11.2 Cached Artifacts

Table 10: Model Cache Contents

File	Contents
<code>propensity_model.pkl</code>	Trained LogisticRegression for $\hat{e}(X)$
<code>outcome_model.pkl</code>	Trained GradientBoostingRegressor for $\hat{l}(X)$
<code>scaler.pkl</code>	Fitted StandardScaler for feature normalization
<code>model_metadata.json</code>	Training metadata (sample count, timestamp, config)

11.3 Progressive Learning Workflow

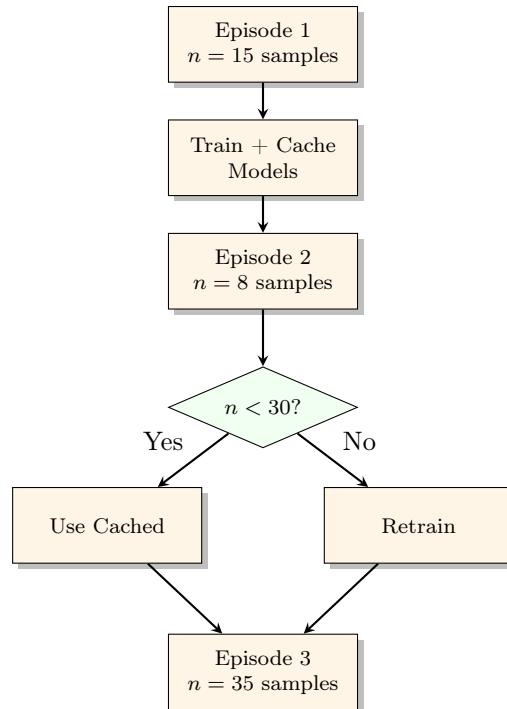


Figure 4: Progressive Learning Across Episodes

12 Integration with RL System

12.1 Usage Example

Listing 3: Basic Usage

```

1 from causal_reward_engine import CausalRewardEngine, CausalConfig
2
3 # Configure engine
4 config = CausalConfig(
5     graph_data_dir="causal_data/graphs",
6     action_file_path="approved_changes.csv",
7     min_graphs_required=10,
8     max_graphs_to_load=20
9 )
10
11 # Create engine
12 engine = CausalRewardEngine(config=config)
13
14 # Compute causal reward
15 result = engine.compute_causal_reward_from_graphs()
16
17 # Extract reward for RL agent
18 causal_reward = result.causal_reward
19 print(f"Causal Reward: {causal_reward:.4f}")
20 print(f"Is Significant: {result.is_significant}")
21 print(f"95% CI: [{result.confidence_interval[0]:.4f}, {result.confidence_interval[1]:.4f}]")

```

12.2 Batch Integration

For offline RL training, replace GNN predictions with causal rewards:

Listing 4: Batch Replacement

```

1 # Traditional approach (using GNN predictions)
2 gnn_reward = gnn_model.predict(state, action)
3
4 # Causal approach (using real data)
5 causal_result = engine.compute_causal_reward_from_graphs(
6     graph_data_dir="episode_graphs/",
7     action_file_path="episode_actions.csv"
8 )
9 causal_reward = causal_result.causal_reward
10
11 # Use causal_reward in experience replay buffer
12 buffer.add(state, action, causal_reward, next_state)

```

13 Computational Complexity

Table 11: Computational Complexity Analysis

Component	Time Complexity	Space Complexity
Load Graphs	$O(G \cdot N)$	$O(G \cdot N)$
QoE Computation	$O(N)$	$O(N)$
PSM	$O(n_1 n_0 \log n_0)$	$O(n_1 + n_0)$
DML (K-fold)	$O(K \cdot n \cdot d \log n)$	$O(n \cdot d)$
Balance Score	$O(p \cdot n)$	$O(p)$
Total	$O(G \cdot N + K \cdot n \cdot d \log n)$	$O(G \cdot N)$

Legend:

- G : Number of graph snapshots (typically 10-20)
- N : Number of APs per graph (typically 5-50)
- $n = n_1 + n_0$: Total samples (treatment + control)
- d : Dimensionality of covariates (typically 7-15)
- K : Number of cross-validation folds (3)
- p : Number of covariates to check for balance

Typical Runtime: For $G = 15$, $N = 20$, $n = 300$, $d = 10$:

- Load + QoE: ~ 0.5 seconds
- PSM: $\sim 1 - 2$ seconds
- DML: $\sim 3 - 5$ seconds (dominated by GBM training)
- **Total:** $\sim 5 - 8$ seconds per evaluation

14 Summary and Conclusion

This document presents a comprehensive design for a Causal Reward Engine that validates RRM actions through rigorous causal inference. The system:

14.1 Key Achievements

- Replaces **correlation with causation** by using PSM and DML to isolate true treatment effects
- Maintains **RL compatibility** by producing rewards in the same format as existing GNN predictions
- Handles **small data** through progressive learning and model persistence
- Provides **statistical guarantees** via significance testing, confidence intervals, and balance checks
- Implements **robust fallbacks** to ensure reliability in production environments

14.2 System Capabilities

- Propensity Score Matching with K-nearest neighbors and caliper constraints
- Double Machine Learning with cross-fitting to prevent overfitting
- Meta-aggregation for robust effect estimation
- Comprehensive balance assessment using Standardized Mean Differences
- Progressive learning with model caching for small sample scenarios
- Sub-10 second inference for real-time deployment

14.3 Integration Benefits

The engine is designed as a drop-in replacement for GNN-based reward prediction, enabling causally-validated RL training for WiFi network optimization.

Key Takeaway: By filtering RL rewards through causal validation, we ensure the agent learns from real improvements rather than spurious correlations, leading to more reliable and effective RRM policies.

14.4 Future Enhancements

- Network interference modeling using graph structure
- Time-series methods with Difference-in-Differences
- Synthetic control for single-unit treatments
- Heterogeneous treatment effect estimation
- Active learning integration with RL exploration