



# **INTER IIT TECH MEET 14.0**

Client Centric WiFi RRM

INTER IIT TECH MEET 14.0

RRM-Plus Complete Architecture

Team 33

**ARISTA**

# Contents

<b>1 Executive Summary</b>	<b>2</b>
<b>2 System Overview</b>	<b>2</b>
2.1 Architecture Philosophy . . . . .	2
<b>3 Component Architecture</b>	<b>2</b>
3.1 NS-3 Simulation Layer . . . . .	2
3.1.1 Time-Based Graph Export . . . . .	3
3.1.2 Multi-Timescale Control Implementation . . . . .	3
3.2 Preprocessing Pipeline . . . . .	3
3.2.1 Noise Floor Modeling . . . . .	3
3.3 Graph Neural Network (GNN) Model . . . . .	4
3.4 Safe Reinforcement Learning Agent . . . . .	4
3.4.1 Action Space . . . . .	5
3.4.2 CQL Loss Function . . . . .	5
3.4.3 Safety Constraint . . . . .	5
3.5 Dream-Based Offline RL . . . . .	5
3.6 Graph Coloring Integration . . . . .	6
3.7 Change Planner with Production Guardrails . . . . .	7
<b>4 Causal Reward Validation: The Key Innovation</b>	<b>7</b>
4.1 Causal Framework . . . . .	8
4.2 Difference-in-Differences Estimator . . . . .	8
4.3 Graph Export Timeline for Causal Validation . . . . .	8
4.4 Causal Reward Computation Algorithm . . . . .	9
<b>5 End-to-End Pipeline</b>	<b>10</b>
5.1 File Dependencies . . . . .	10
<b>6 Performance Analysis</b>	<b>10</b>
6.1 Computational Complexity . . . . .	10
6.2 Scalability Analysis . . . . .	11
<b>7 Deployment Architecture</b>	<b>12</b>
7.1 System Components . . . . .	12
7.2 API Design . . . . .	12
<b>8 Conclusion and Future Work</b>	<b>13</b>
8.1 Key Contributions . . . . .	13
8.2 Future Enhancements . . . . .	13
8.3 Impact . . . . .	13
<b>A Appendix: Mathematical Formulations</b>	<b>13</b>
A.1 GNN Message Passing . . . . .	13
A.2 Conservative Q-Learning . . . . .	14
A.3 Difference-in-Differences with Fixed Effects . . . . .	14
<b>B Appendix: File Structure</b>	<b>14</b>

# 1 Executive Summary

This document presents the comprehensive architecture design for **RRM-Plus**, an AI-assisted, client-centric Radio Resource Management system developed for Arista Networks' enterprise WiFi infrastructure. The system represents a paradigm shift from traditional AP-centric RRM approaches by leveraging:

- **Graph Neural Networks (GNNs)** for interference topology learning
- **Safe Reinforcement Learning** with Conservative Q-Learning and production guardrails
- **Multi-timescale control loops** (fast/slow/event-driven)
- **Causal inference** for reward validation against counterfactual holdouts
- **NS-3 simulation** with realistic 5GHz noise modeling

**Key Innovation:** Unlike traditional RL systems that rely on simple delta-QoE rewards, RRM-Plus employs **causal reward validation** through difference-in-differences estimation. This ensures that observed QoE improvements are causally attributable to RL actions rather than confounding factors (time-of-day effects, client mobility, external interference).

## Performance Targets:

- +25-35% median throughput for edge clients
- -30% P95 retry rate network-wide
- <0.2 configuration changes per AP per day (stability)
- >90% BSS-TM steering acceptance

# 2 System Overview

## 2.1 Architecture Philosophy

RRM-Plus adopts a **hybrid control architecture** that combines:

1. **Physics-based simulation** (NS-3) for realistic network dynamics
2. **Graph-structured learning** (GNN) for spatial interference relationships
3. **Safe offline RL** (Conservative Q-Learning) for policy optimization
4. **Production guardrails** (change budgets, blast-radius control)
5. **Causal validation** (difference-in-differences, propensity matching)

This architecture describes both the NS-3 simulation pipeline and its intended real-world deployment counterpart. Sections 3–6 refer to the simulation system, while Section 7 outlines the production-level architecture.

# 3 Component Architecture

## 3.1 NS-3 Simulation Layer

**File:** basic-rrm-phase23\_9th\_nov.cc

The NS-3 simulation forms the foundation of RRM-Plus, providing:

- **802.11ax PHY/MAC simulation** with OFDMA, MU-MIMO, spatial reuse
- **Realistic propagation models** (ITU Indoor, log-distance)
- **5GHz noise floor simulation** with temporal dynamics
- **Client mobility and traffic patterns**
- **Interference graph export** (adjacency + coupling strengths)

### 3.1.1 Time-Based Graph Export

Every 30 seconds, NS-3 invokes the Python controller:

```

1 bool DataGenerationController::InvokeGraphColoring(...) {
2     Time now = Simulator::Now();
3     double currentTimeSeconds = now.GetSeconds();
4     int currentTimeInt = static_cast<int>(std::round(currentTimeSeconds));
5
6     // Enable causal analysis every 30s
7     bool enableCausal = (currentTimeInt % 30 == 0);
8
9     cmd << " --nodes " << nodesFile;
10    cmd << " --edges " << edgesFile;
11    cmd << " --flag-causal " << (enableCausal ? "true" : "false");
12    cmd << " --num-dreams 5";
13
14    system(cmd.str().c_str());
15 }
```

### 3.1.2 Multi-Timescale Control Implementation

Table 1: Control Loop Timescales

Loop Type	Period	Triggers	Actions
Fast	5-15s	DFS events, Noise spikes	Channel change, OBSS-PD
Slow	30-60s	AI optimization	Power, Width, Channel
Event	Async	Microwave, BLE bursts	Transient mitigation

## 3.2 Preprocessing Pipeline

File: `preprocessing.py`

Transforms raw CSV exports into PyTorch Geometric Data objects:




Figure 2: Preprocessing pipeline with 5GHz noise modeling

### 3.2.1 Noise Floor Modeling

The `noise.py` module simulates realistic 5GHz office environments:

$$\eta(t, \text{ap\_id}) = \eta_{\text{base}}(t) + \eta_{\text{lunch}}(t) + \eta_{\text{AV}}(t) + \eta_{\text{AP\_bias}} + \xi(t) \quad (1)$$

$$\eta_{\text{base}}(t) = -90 \text{ dBm} \quad (\text{workday plateau}) \quad (2)$$

$$\eta_{\text{lunch}}(t) = 5 \cdot \mathcal{N}(t | \mu = 12.5, \sigma = 0.4) \quad (\text{microwave spikes}) \quad (3)$$

$$\eta_{\text{AV}}(t) \sim \text{Bernoulli}(p = 0.2) \cdot \mathcal{U}(5, 8) \text{ dBm} \quad (\text{wireless casting}) \quad (4)$$

### 3.3 Graph Neural Network (GNN) Model

File: `model.py`

Architecture: **GraphSAGE** with attention-based aggregation




Figure 3: GNN architecture for learning interference topology

#### Key Design Choices:

- **GraphSAGE** over GCN: better generalization to unseen topologies
- **Multi-layer aggregation:** captures 3-hop interference relationships
- **Dual prediction heads:** simultaneous QoE and safety estimation
- **Residual connections:** prevent gradient vanishing in deep graphs

### 3.4 Safe Reinforcement Learning Agent

File: `agent_v2.py`

Implementation: **Conservative Q-Learning (CQL)** with Lagrangian safety constraints

### 3.4.1 Action Space

Discrete action space:  $\mathcal{A} = \mathcal{A}_{\text{power}} \times \mathcal{A}_{\text{OBSS}} \times \mathcal{A}_{\text{width}}$

- $\mathcal{A}_{\text{power}} = \{1, 3, 5, \dots, 30\}$  dBm (15 values)
- $\mathcal{A}_{\text{OBSS}} = \{-92, -82, -72, -62, -52\}$  dBm (5 values)
- $\mathcal{A}_{\text{width}} = \{20, 40, 80\}$  MHz (3 values)

Total actions:  $|\mathcal{A}| = 15 \times 5 \times 3 = 225$

### 3.4.2 CQL Loss Function

$$\mathcal{L}_{\text{CQL}}(\theta) = \underbrace{\alpha \cdot \mathbb{E}_{s \sim \mathcal{D}} \left[ \log \sum_a \exp Q_\theta(s, a) - \mathbb{E}_{a \sim \pi_\beta} [Q_\theta(s, a)] \right]}_{\text{Pessimism penalty}} \quad (5)$$

$$+ \underbrace{\mathbb{E}_{(s, a, r, s') \sim \mathcal{D}} \left[ (Q_\theta(s, a) - (r + \gamma \max_{a'} Q_{\theta'}(s', a')))^2 \right]}_{\text{Bellman error}} \quad (6)$$

where:

- $\alpha$ : CQL coefficient (controls pessimism level)
- $\pi_\beta$ : behavior policy (offline dataset)
- $\mathcal{D}$ : offline replay buffer

### 3.4.3 Safety Constraint

Lagrangian formulation for safety-aware RL:

$$\max_{\pi} \mathbb{E}_\pi [R(s, a)] \quad (7)$$

$$\text{s.t. } \mathbb{E}_\pi [C(s, a)] \leq d_{\text{safety}} \quad (8)$$

Implemented via dual gradient ascent on  $\lambda$  (safety multiplier):

$C(s, a)$  represents predicted safety violations (e.g., airtime > threshold, latency > limit), as predicted by the Safety Head.

```

1 # Update Q-functions
2 q_r_loss = cql_loss_reward + bellman_loss_reward
3 q_c_loss = cql_loss_cost + bellman_loss_cost
4
5 # Update safety multiplier (lambda)
6 costViolation = q_c_pred.mean() - self.safety_threshold
7 lambda_loss = -self.lambda_param * costViolation
8 lambda_new = max(0.0, self.lambda_param + self.lambda_lr * costViolation)

```

## 3.5 Dream-Based Offline RL

**Key Innovation:** Generate synthetic transitions using the GNN as a world model

**Algorithm 1** Dream Sample Generation

---

```

1: Input: Current graph  $G_t$ , GNN model  $f_\theta$ , action  $a$ 
2: Output: Dream transition  $(s_t, a, r, s_{t+1})$ 
3:
4:  $\mathbf{h}_t \leftarrow f_\theta(G_t)$                                  $\triangleright$  GNN embedding
5:  $\hat{y}_{\text{current}} \leftarrow \text{QoE\_head}(\mathbf{h}_t)$            $\triangleright$  Current KPIs
6:
7:  $G_{\text{dream}} \leftarrow \text{ApplyAction}(G_t, a)$             $\triangleright$  Modify node features
8:  $\mathbf{h}_{\text{dream}} \leftarrow f_\theta(G_{\text{dream}})$ 
9:  $\hat{y}_{\text{dream}} \leftarrow \text{QoE\_head}(\mathbf{h}_{\text{dream}})$            $\triangleright$  Predicted KPIs
10:
11:  $\Delta_{\text{QoE}} \leftarrow \hat{y}_{\text{dream}} - \hat{y}_{\text{current}}$ 
12:  $r \leftarrow w_1 \cdot \Delta_{\text{throughput}} + w_2 \cdot \Delta_{\text{retry}} + w_3 \cdot \Delta_{\text{latency}}$ 
13:
14: safety\_score  $\leftarrow \text{Safety\_head}(\mathbf{h}_{\text{dream}})$ 
15:  $c \leftarrow \mathbb{1}[\text{safety\_score} < \theta_{\text{safe}}]$                    $\triangleright$  Cost signal
16:
17: return  $(s_t, a, r, c, s_{t+1})$ 

```

---

**Advantages:**

- **Data efficiency:** 5 dreams per real transition  $\rightarrow 5\times$  sample efficiency
- **Safety:** test risky actions in simulation before deployment
- **Exploration:** try actions outside offline distribution

### 3.6 Graph Coloring Integration

**File:** `rl_gc_integration.py`

Combines RL power/OBSS-PD decisions with DSatur graph coloring for channels:




Figure 4: RL-Graph Coloring integration for joint optimization

**Why separate RL and Graph Coloring?**

- **Scalability:** RL on full joint space ( $225^N$  actions for  $N$  APs) is intractable
- **Optimality:** DSatur provides near-optimal channel assignment with proven bounds
- **Constraints:** Graph coloring naturally enforces regulatory/DFS constraints

### 3.7 Change Planner with Production Guardrails

**File:** `change_planner.cpp`

Enforces production safety policies before applying RL recommendations:

Table 2: Production Guardrails

Guardrail	Constraint	Rationale
Change Budget	$\leq 1$ change per AP per 4h	Prevent flapping
Blast Radius	$\leq N$ APs per RF domain	Limit failure scope
Hysteresis	Power $\geq 2$ dB, Width $\geq 1$ step	Avoid micro-changes
Time Window	No changes during peak hours	Preserve user experience
Rollback	Auto-revert if KPI $\downarrow X\%$	Safety net

---

#### Algorithm 2 Change Planner Decision Logic

---

```

1: Input: Proposed changes  $\mathcal{C} = \{c_1, \dots, c_N\}$ 
2: Output: Actions {APPLY, SKIP, ROLLBACK}
3:
4: for each change  $c_i$  do
5:   if  $c_i$  violates change budget then
6:      $action_i \leftarrow \text{SKIP}$ 
7:   else if  $c_i$  violates blast radius then
8:      $action_i \leftarrow \text{SKIP}$ 
9:   else if  $\Delta_{\text{param}} <$  hysteresis threshold then
10:     $action_i \leftarrow \text{SKIP}$ 
11:   else if current time in blackout window then
12:      $action_i \leftarrow \text{SKIP}$ 
13:   else
14:      $action_i \leftarrow \text{APPLY}$ 
15:   end if
16: end for
17:
18: Apply changes marked APPLY
19: Monitor KPIs for 30s
20: if KPI degradation  $> \theta_{\text{rollback}}$  then
21:   Revert changes (ROLLBACK)
22: end if

```

---

## 4 Causal Reward Validation: The Key Innovation

**Problem:** Traditional RL rewards based on  $\Delta_{\text{QoE}}$  are confounded by:

- Time-of-day effects (lunch traffic spike)
- Client mobility (users entering/leaving coverage)
- External interference (microwave, BLE)

**Solution:** Use causal inference to estimate the **causal effect** of RL actions

#### 4.1 Causal Framework




Figure 5: Causal graph showing confounding factors

**Goal:** Estimate Average Treatment Effect (ATE):

$$\tau_{ATE} = \mathbb{E}[Y(1) - Y(0)]$$

where  $Y(1)$  is QoE under treatment (apply),  $Y(0)$  is QoE under control (skip).

#### 4.2 Difference-in-Differences Estimator

**File:** `causal_reward_engine.py`

Implementation uses a **difference-in-differences (DiD)** approach:

$$\hat{\tau}_{DiD} = (\bar{Y}_{treated}^{post} - \bar{Y}_{treated}^{pre}) - (\bar{Y}_{control}^{post} - \bar{Y}_{control}^{pre}) \quad (9)$$




Figure 6: Difference-in-differences visualization: causal effect is the gap between treatment and counterfactual

#### 4.3 Graph Export Timeline for Causal Validation

After applying changes, NS-3 exports 6 graph snapshots at 5s intervals:




Figure 7: Timeline of graph exports for causal validation, snapshots G7...G32 capture post-intervention stabilization. The 6-snapshot averaging reduces transient variance.

#### 4.4 Causal Reward Computation Algorithm

---

##### Algorithm 3 Causal Reward Engine

---

```

1: Input: Graph snapshots  $\{G_0, G_7, G_{12}, G_{17}, G_{22}, G_{27}, G_{32}\}$ , action log
2: Output: Causal reward  $r_{\text{causal}}$ 
3:
4: // Step 1: Identify treatment and control groups
5:  $\mathcal{T} \leftarrow \{\text{AP}_i : \text{action}_i = \text{APPLY}\}$ 
6:  $\mathcal{C} \leftarrow \{\text{AP}_i : \text{action}_i = \text{SKIP}\}$ 
7:
8: // Step 2: Extract QoE trajectories
9: for each AP  $i$  do
10:    $Y_i^{\text{pre}} \leftarrow \text{QoE}(G_0, i)$                                      ▷ Baseline
11:    $Y_i^{\text{post}} \leftarrow \frac{1}{6} \sum_{t \in \{7, 12, 17, 22, 27, 32\}} \text{QoE}(G_t, i)$            ▷ Average post-intervention
12: end for
13:
14: // Step 3: Compute difference-in-differences
15:  $\Delta_{\mathcal{T}} \leftarrow \frac{1}{|\mathcal{T}|} \sum_{i \in \mathcal{T}} (Y_i^{\text{post}} - Y_i^{\text{pre}})$ 
16:  $\Delta_{\mathcal{C}} \leftarrow \frac{1}{|\mathcal{C}|} \sum_{i \in \mathcal{C}} (Y_i^{\text{post}} - Y_i^{\text{pre}})$ 
17:
18:  $\hat{\tau}_{\text{DiD}} \leftarrow \Delta_{\mathcal{T}} - \Delta_{\mathcal{C}}$                                          ▷ Causal effect
19:
20: // Step 4: Statistical significance test
21:  $p \leftarrow \text{WelchTest}(\{Y_i^{\text{post}} - Y_i^{\text{pre}}\}_{i \in \mathcal{T}}, \{Y_i^{\text{post}} - Y_i^{\text{pre}}\}_{i \in \mathcal{C}})$ 
22: if  $p < 0.05$  then
23:    $r_{\text{causal}} \leftarrow \hat{\tau}_{\text{DiD}}$                                          ▷ Statistically significant uplift
24: else
25:    $r_{\text{causal}} \leftarrow 0$                                               ▷ No significant effect
26: end if
27:
28: return  $r_{\text{causal}}$ 

```

---

#### Key Features:

- **Treatment/control separation:** APPLY vs SKIP actions create natural A/B test
- **Parallel trends assumption:** Both groups experience same time-of-day effects
- **Statistical testing:** Welch's t-test ensures effects are not due to chance
- **Robustness:** 6 post-intervention snapshots provide stable estimate

## 5 End-to-End Pipeline

### 5.1 File Dependencies




Figure 8: Broad file dependency graph

## 6 Performance Analysis

### 6.1 Computational Complexity

Table 3: Time Complexity Analysis

Component	Complexity	Typical Time
NS-3 Simulation (1s)	$O(N \cdot M \cdot T)$	0.5-1.0s
Graph Export	$O(N + E)$	0.1s
Preprocessing	$O(N + E)$	0.05s
GNN Forward Pass	$O(L \cdot E \cdot D)$	0.02s
RL Inference	$O(N \cdot A)$	0.01s
Graph Coloring	$O(N^2)$	0.1s
Change Planning	$O(N)$	0.01s
Causal Analysis	$O(6N)$	2-5s
<b>Total (no causal)</b>		<b>0.8s</b>
<b>Total (with causal)</b>		<b>3-5s</b>

where:

- $N$ : number of APs (typically 5-50)
- $M$ : number of clients (typically 10-500)
- $E$ : number of edges in interference graph ( $O(N^2)$ )

- $L$ : GNN layers (3)
- $D$ : embedding dimension (256)
- $A$ : action space size (225)

## 6.2 Scalability Analysis




Figure 9: Scalability of pipeline components

### Key Insights:

- GNN scales linearly with  $E$  (message passing)
- RL scales linearly with  $N$  (independent per-AP decisions)
- Graph coloring is the bottleneck at large scale (quadratic)

### Optimization:

- **Sparse graphs:** Prune edges with coupling < -90 dBm
- **Hierarchical coloring:** Decompose into RF domains
- **Approximate algorithms:** Greedy coloring with  $O(N \log N)$

## 7 Deployment Architecture

### 7.1 System Components




Figure 10: Production deployment architecture

### 7.2 API Design

#### Key APIs:

```

1 # /api/v1/inference
2 POST /api/v1/inference
3 {
4     "graph_nodes": [...],
5     "graph_edges": [...],
6     "current_config": [...],
7     "timestamp": "2025-12-05T14:30:00Z"
8 }
9 Response: {
10     "proposed_changes": [...],
11     "predicted_qoe_delta": 12.5,
12     "confidence": 0.87
13 }
14
15 # /api/v1/causal/validate
16 POST /api/v1/causal/validate
17 {
18     "graph_snapshots": [...],
19     "action_log": [...],
20     "method": "did"
21 }
22 Response: {
23     "causal_effect": 18.3,
24     "p_value": 0.012,
25     "confidence_interval": [12.1, 24.5]
26 }
```

## 8 Conclusion and Future Work

### 8.1 Key Contributions

1. **Client-centric RRM:** First system to optimize for client QoE rather than AP metrics
2. **Safe offline RL:** Conservative Q-Learning with production guardrails for deployment
3. **Graph-structured learning:** GNN captures spatial interference relationships
4. **Causal validation:** DiD estimator eliminates confounding for accurate credit assignment
5. **Multi-timescale control:** Unified framework for fast/slow/event-driven loops

### 8.2 Future Enhancements

- **Online learning:** Continual adaptation to changing environments
- **Multi-site transfer:** Pre-trained models for rapid deployment
- **Federated learning:** Privacy-preserving collaboration across sites
- **Causal discovery:** Auto-learn causal graph structure (PC algorithm)
- **6 GHz support:** Extend to WiFi 6E with AFC integration
- **Explainable AI:** SHAP values for interpretable decisions

### 8.3 Impact

RRM-Plus demonstrates that **AI + Physics + Causality** can deliver:

- QoE improvements in dense deployments
- Safe, stable, production-ready automation
- Validated performance via causal inference

This will enable **intelligent, self-optimizing WiFi infrastructure** for enterprise and campus networks.

## A Appendix: Mathematical Formulations

### A.1 GNN Message Passing

$$\mathbf{m}_{i \leftarrow j}^{(l)} = \phi_{\text{msg}}(\mathbf{h}_i^{(l)}, \mathbf{h}_j^{(l)}, \mathbf{e}_{ij}) \quad (10)$$

$$\mathbf{h}_i^{(l+1)} = \phi_{\text{update}} \left( \mathbf{h}_i^{(l)}, \bigoplus_{j \in \mathcal{N}(i)} \mathbf{m}_{i \leftarrow j}^{(l)} \right) \quad (11)$$

For GraphSAGE:

$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{W}^{(l)} \cdot \text{CONCAT} \left( \mathbf{h}_i^{(l)}, \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} \mathbf{h}_j^{(l)} \right) \right) \quad (12)$$

## A.2 Conservative Q-Learning

Full objective with Lagrangian safety:

$$\mathcal{L}(\theta) = \alpha \left[ \mathbb{E}_{s \sim \mathcal{D}} \left[ \log \sum_a \exp Q_\theta(s, a) \right] - \mathbb{E}_{s, a \sim \mathcal{D}} [Q_\theta(s, a)] \right] \quad (13)$$

$$+ \mathbb{E}_{(s, a, r, s') \sim \mathcal{D}} [(Q_\theta(s, a) - \hat{Q}(s, a))^2] \quad (14)$$

$$+ \lambda \cdot \mathbb{E}_{s \sim \mathcal{D}} \left[ \max_a Q_\theta^C(s, a) - d_{\text{safety}} \right] \quad (15)$$

where  $\hat{Q}(s, a) = r + \gamma \max_{a'} Q_\theta(s', a')$  is the Bellman target.

## A.3 Difference-in-Differences with Fixed Effects

Panel data regression:

$$Y_{it} = \beta_0 + \beta_1 \cdot \text{Treat}_i + \beta_2 \cdot \text{Post}_t + \beta_3 \cdot (\text{Treat}_i \times \text{Post}_t) \quad (16)$$

$$+ \alpha_i + \gamma_t + \epsilon_{it} \quad (17)$$

where:

- $\alpha_i$ : AP fixed effects (time-invariant)
- $\gamma_t$ : Time fixed effects (common trends)
- $\beta_3$ : Causal effect (DiD estimator)

## B Appendix: File Structure

```

project/
  ns-3/
    scratch/
      basic-rrm-phase23_9th_nov.cc
      change_planner.cpp
  python/
    python_controller.py
    rl_gc_integration.py
    agent_v2.py
    model.py
    arista_dataset_v2.py
    preprocessing.py
    noise.py
    utils.py
    causal_reward_engine.py
    explainability.py
    gc_py.py
    json_export.py
  logs/
    interference_graph_nodes/
    interference_graph_edges/
    RL_GNN/
      RL_Action/
        Nodes/
        Edges/
        export_status.txt

```

```
    causal_analysis/  
data/  
    approved_changes.csv  
    current_config.csv  
    rl_suggestions_copy.csv  
docs/  
    CAUSAL_RL_INTEGRATION.md  
    architecture.tex
```